

Tajamul Khan

+91 1234567890 | India | [me@gmail.com](#) | [LinkedIn](#) | [Github](#) | [Portfolio](#)

SUMMARY

Senior Data Scientist & AI Engineer with 7+ years of experience in machine learning, Generative AI, and Agentic AI. Specialized in production-grade RAG pipelines, multi-agent systems, and LLM-powered applications with proven impact — \$5M quarterly savings, President's Award at Google, and 3 Microsoft certifications.

EXPERIENCE

Advanced Data Scientist

Eastman – Global Capability Center

July 2000 – Present

Hyderabad, India

- Lead production AI initiatives: RAG-based document intelligence, LLM fine-tuning (PEFT), customer segmentation, and demand forecasting on Azure. Collaborate with cross-functional teams to translate AI outputs into strategic business decisions.

Data Scientist - Product

Google

May 2000 – June 2000

Hyderabad, India

- Delivered end-to-end analytics — competitive analysis, sales forecasting, user behavior modeling. Received the **President's Award** for data-led strategies that enhanced operational efficiency and revenue impact.

Data Quality Specialist

Amazon Development Centre

Nov 2000 – Oct 2000

Hyderabad, India

- Automated data preprocessing and reporting pipelines using Python and SQL; built interactive dashboards to visualize KPIs and support data-driven decisions.

EDUCATION

Massachusetts Institute of Technology – MIT

Post Graduate Program – Data Science and Machine Learning

May 2000 – Nov 2000

Kurukshetra University

Master of Engineering

Aug 2000 – May 2000

PROJECTS

- **Enterprise RAG Copilot for Call Reports** *Azure OpenAI · GPT-4o Mini · HNSW · PEFT · Azure AI Search*
Built a production RAG system for PDF-based call report intelligence. Fine-tuned GPT-4o Mini with PEFT; implemented HNSW vector search across 10K+ docs. Evaluated via Hit Rate, MRR, and semantic relevancy — outperforming humans on 105/300+ queries. **Reduced manual effort 50%, saving \$5M/quarter.**
- **Multi-Agent Market Intelligence System** *LangGraph · CrewAI · Python · SQL · Power BI*
Designed a multi-agent system with LangGraph and CrewAI: specialized agents for scraping, SQL enrichment, trend analysis, and report generation — orchestrated by a supervisor agent with human-in-the-loop gates. **Cut research turnaround from 3 days to 2 hours**, driving significant revenue pipeline growth.
- **Real-Time Fraud Detection Platform** *MLflow · Docker · FastAPI · GitHub Actions · Azure ML · Evidently AI*
Engineered an end-to-end fraud detection system with automated retraining triggered by data drift alerts via Evidently AI. Built ensemble models (XGBoost + Isolation Forest) with sub-100ms inference via FastAPI; containerized with Docker and deployed on Azure ML. CI/CD via GitHub Actions ensured zero-downtime releases. **Reduced false positives by 35% and cut deployment cycle from 2 weeks to 1 day.**

SKILLS

- **Languages & Tools:** Python, SQL, R, PySpark, Power BI, Tableau, Streamlit, Flask
- **AI & ML:** RAG Pipelines, Multi-Agent Systems (LangGraph, CrewAI), LLM Fine-Tuning (PEFT/LoRA), NLP, Deep Learning, Time Series Forecasting
- **Cloud & MLOps:** Azure OpenAI, Azure ML, MLflow, Docker, GitHub Actions, Vector Databases (HNSW, FAISS), Microsoft Fabric

CERTIFICATIONS

Microsoft Certified Fabric Analytics Engineer (DP-600): [Link](#) Microsoft Certified Fabric Data Engineer (DP-700): [Link](#)
Microsoft Certified Power BI Data Analyst (PL-300): [Link](#) Google Certified Data Analyst: [Link](#)